

WEBVTT

1

00:00:00.900 --> 00:00:03.700

<v ->What happens when you have greater cloud choice?</v>

2

00:00:03.700 --> 00:00:07.310

With VMware cloud on AWS, all your apps can benefit

3

00:00:07.310 --> 00:00:11.340

from a cloud that gives you choice, flexibility, and speed.

4

00:00:11.340 --> 00:00:13.703

VMware, welcome change.

5

00:00:15.258 --> 00:00:18.675 line:15%

(relaxing music playing)

6

00:00:25.540 --> 00:00:29.840

<v ->Hi, welcome to beyond five 9s,</v>

7

00:00:29.840 --> 00:00:31.790

which is all about lessons from our highest

8

00:00:31.790 --> 00:00:33.170

available data planes.

9

00:00:33.170 --> 00:00:35.920

I'm Colm, I'm a VP and distinguished engineer here

10

00:00:35.920 --> 00:00:39.704

at Amazon web services, we're gonna be joined by Yasemin,

11

00:00:39.704 --> 00:00:41.950

who's a principal engineer at Amazon web services,

12

00:00:41.950 --> 00:00:43.800

she works on Kinesis.

13

00:00:43.800 --> 00:00:48.640

And today we're gonna share a bunch of insights into

14

00:00:48.640 --> 00:00:50.990

how we build our most highly available systems.

15

00:00:50.990 --> 00:00:53.330

You know, what's going on under the hood

16

00:00:53.330 --> 00:00:55.113

or behind the scenes, so to speak.

17

00:00:56.210 --> 00:00:58.170

You know, we've gone into great detail about

18

00:00:58.170 --> 00:01:01.690

some of these topics in our Amazon builders library series.

19

00:01:01.690 --> 00:01:05.310

But today we're gonna cherry pick some of the best

20

00:01:06.320 --> 00:01:08.880

lessons and condense them down into, you know,

21

00:01:08.880 --> 00:01:12.400

10 simple things hopefully you can take away and apply

22

00:01:12.400 --> 00:01:15.040

to your own systems as you build them.

23

00:01:15.040 --> 00:01:19.285

These, these lessons come from systems that have, you know,

24

00:01:19.285 --> 00:01:23.500

stood the test of time and had very good availability

25

00:01:23.500 --> 00:01:25.733

records, which is not easy to achieve.

26

00:01:27.290 --> 00:01:29.220

We've talked about this before,

27

00:01:29.220 --> 00:01:32.274

and some of the lessons are eternal and,

28

00:01:32.274 --> 00:01:33.860

and some of the lessons are new.

29

00:01:33.860 --> 00:01:35.730

But the first one that we're gonna start with is,

30

00:01:35.730 --> 00:01:39.370

is definitely one that's eternal, which is it's,

31

00:01:39.370 --> 00:01:41.840

it's all about insisting on high standards, right?

32

00:01:41.840 --> 00:01:45.184

Nothing we're gonna go into, there's no tip or trick

33

00:01:45.184 --> 00:01:48.740

or magic that we have for you that can ever replace

34

00:01:48.740 --> 00:01:52.600

paying very strong attention to detail in,

35

00:01:52.600 --> 00:01:54.743

in how we build systems.

36

00:01:56.260 --> 00:01:59.080

There's a few things I'm always reminded about

37

00:01:59.080 --> 00:02:01.040

when I think of this.

38

00:02:01.040 --> 00:02:02.870

One is something I learned from Werner Vogels,

39

00:02:02.870 --> 00:02:05.330

our CTO a long time ago, which is when,

40

00:02:05.330 --> 00:02:07.130

which is that when you have, you know,

41

00:02:07.130 --> 00:02:09.830

systems that are processing billions of requests,

42

00:02:09.830 --> 00:02:12.283

or you know, even trillions of requests over,

43

00:02:13.165 --> 00:02:15.850

you know, short periods of time,

44

00:02:15.850 --> 00:02:18.130

pretty much anything that could go wrong in the system

45

00:02:18.130 --> 00:02:21.620

will be going wrong at some point in time,

46

00:02:21.620 --> 00:02:25.110

some request somewhere is probably experiencing, you know,

47

00:02:25.110 --> 00:02:26.940

even a one in a billion failure, right?

48

00:02:26.940 --> 00:02:28.950

Because you have billions of requests.

49

00:02:28.950 --> 00:02:32.940

And so it requires paying enormous attention to detail,

50

00:02:32.940 --> 00:02:34.800

looking for every single error case,

51

00:02:34.800 --> 00:02:37.005

looking for everything that could possibly go wrong

52

00:02:37.005 --> 00:02:39.983

and having a plan for that in the code and in testing

53

00:02:39.983 --> 00:02:44.670

and everywhere, and that kind of culture, at team level,

54

00:02:44.670 --> 00:02:46.710

there's just no substitute for it.

55

00:02:46.710 --> 00:02:50.710

You know, having teams who are willing to constantly raise

56

00:02:50.710 --> 00:02:53.760

the bar on testing, constantly, you know, write new

57

00:02:53.760 --> 00:02:56.780

and interesting tests or come up with ways to make

58

00:02:56.780 --> 00:02:58.330

testing itself easier.

59

00:02:58.330 --> 00:03:01.900

You know, a lot of that is where the quality

60

00:03:01.900 --> 00:03:03.837

and high availability in our systems comes from,

61

00:03:03.837 --> 00:03:07.310

and like I say, there's no substitute for it.

62

00:03:07.310 --> 00:03:11.730

And I often think that that kind of quality comes about

63

00:03:11.730 --> 00:03:13.060
from great habits, right?

64

00:03:13.060 --> 00:03:16.090
It's from what we do over and over and over again.

65

00:03:16.090 --> 00:03:19.177
And it's really hard to approach a system and go, you know,

66

00:03:19.177 --> 00:03:20.870
"I'm just gonna add quality to it.

67

00:03:20.870 --> 00:03:23.890
I'm just gonna try to make it better than it was before."

68

00:03:23.890 --> 00:03:26.090
It's more like something pervasive that has to be there

69

00:03:26.090 --> 00:03:27.500
from the beginning.

70

00:03:27.500 --> 00:03:30.470
And so something we've been paying special attention to

71

00:03:30.470 --> 00:03:32.780
on this kind of a quality dimension for,

72

00:03:32.780 --> 00:03:37.780
for a long time now is deployment safety, where,

73

00:03:37.840 --> 00:03:41.560
you know, we write our code and we go through code review,

74

00:03:41.560 --> 00:03:42.660
collaborative code review,

75

00:03:42.660 --> 00:03:46.410

and so on and back and forth processes like you'd be very

76

00:03:46.410 --> 00:03:49.470

familiar with, but once we've checked in our code,

77

00:03:49.470 --> 00:03:52.530

we actually have quite sophisticated deployment systems

78

00:03:52.530 --> 00:03:55.900

that take over after that and kind of take charge

79

00:03:55.900 --> 00:03:57.930

of getting that code from, you know,

80

00:03:57.930 --> 00:03:59.370

the fact that it's been checked in,

81

00:03:59.370 --> 00:04:01.690

all the way to running in production.

82

00:04:01.690 --> 00:04:05.030

And that system has been designed in a way where it

83

00:04:05.030 --> 00:04:09.000

generates a lot of positive safety for our systems.

84

00:04:09.000 --> 00:04:11.650

You know, whenever we're deploying new code, right?

85

00:04:11.650 --> 00:04:13.530

That always means some risks.

86

00:04:13.530 --> 00:04:16.740

We try to be very paranoid about our testing.

87

00:04:16.740 --> 00:04:18.210

We try to be very creative in our testing

88

00:04:18.210 --> 00:04:21.230

and test it in every possible way, and I'll talk more

89

00:04:21.230 --> 00:04:23.570

about that later, but you know,

90

00:04:23.570 --> 00:04:26.030

new code that hasn't been run in production before,

91

00:04:26.030 --> 00:04:28.090

there's always some chance that it might encounter

92

00:04:28.090 --> 00:04:30.910

some condition that we didn't plan for.

93

00:04:30.910 --> 00:04:35.430

And so we've got, you know, CI/CD processes that make sure

94

00:04:35.430 --> 00:04:38.410

that, you know, without us having to be too involved,

95

00:04:38.410 --> 00:04:40.950

that code will go through, you know,

96

00:04:40.950 --> 00:04:44.870

pretty careful, deliberate testing and deployment

97

00:04:44.870 --> 00:04:46.050

as it goes, right?

98

00:04:46.050 --> 00:04:48.090

So we'll deploy it to, you know,

99

00:04:48.090 --> 00:04:52.040

maybe just one box to start with and see how it goes,

100

00:04:52.040 --> 00:04:55.200

there'll be some tests run and that box will have to pass

101

00:04:55.200 --> 00:04:57.790

all of those tests in every condition.

102

00:04:57.790 --> 00:04:58.870

And if it doesn't, you know,

103

00:04:58.870 --> 00:05:01.270

we've got fast and reliable rollback,

104

00:05:01.270 --> 00:05:03.640

that's obviously pretty important too,

105

00:05:03.640 --> 00:05:06.110

but if it does go well, you know, it gets promoted.

106

00:05:06.110 --> 00:05:07.580

And the idea is, you know,

107

00:05:07.580 --> 00:05:10.790

the more and more confidence we have in that code,

108

00:05:10.790 --> 00:05:15.710

the broader we're, we're willing to run that code, right?

109

00:05:15.710 --> 00:05:18.410

So, as we develop confidence in it because it's passing

110

00:05:18.410 --> 00:05:21.790

more and more tests, we'll promote it from just one box

111

00:05:21.790 --> 00:05:26.520

to one cell maybe, or one availability zone, one region,

112

00:05:26.520 --> 00:05:28.540

multiple regions and so on.

113

00:05:28.540 --> 00:05:31.080

And this is something, you know, we've gone really,

114

00:05:31.080 --> 00:05:32.320

really deep at our culture.

115

00:05:32.320 --> 00:05:35.190

You know, every development team at AWS knows

116

00:05:35.190 --> 00:05:39.440

and understands this deployment process pretty intricately

117

00:05:39.440 --> 00:05:42.010

and knows how to, you know, get the most out of it

118

00:05:42.010 --> 00:05:45.950

and design their deployments and code in a way where

119

00:05:45.950 --> 00:05:47.940

they're gonna extract the most benefit from it,

120

00:05:47.940 --> 00:05:49.800

and We've seen a lot of return from this.

121

00:05:49.800 --> 00:05:52.940

It's been, it's been, you know,

122

00:05:52.940 --> 00:05:56.430

a great mechanism for us to deliver

123

00:05:56.430 --> 00:05:58.500

high availability in systems.

124

00:05:58.500 --> 00:06:02.060

And, you know, the occasional issues that it catches

125

00:06:02.060 --> 00:06:03.490

are well, well worth it.

126

00:06:03.490 --> 00:06:04.890

So that's, that's our first,

127

00:06:05.780 --> 00:06:07.590

first thing, and the first example of insisting

128

00:06:07.590 --> 00:06:09.140

on high standards.

129

00:06:09.140 --> 00:06:11.810

The next, and something we haven't really talked too much

130

00:06:11.810 --> 00:06:16.270

about before is, is how we manage our systems, right?

131

00:06:16.270 --> 00:06:18.450

So there's, there's this concept out there, which,

132

00:06:18.450 --> 00:06:20.687

which really condensed it down for me, and I love about,

133

00:06:20.687 --> 00:06:23.850

you know, you want to manage your systems,

134

00:06:23.850 --> 00:06:26.780

more like cattle than like kept pets in that you want to be

135

00:06:26.780 --> 00:06:31.250

able to think of them in abstractions and not, you know, be,

136

00:06:31.250 --> 00:06:34.410

be managing, you know, literally one host or one server

137

00:06:34.410 --> 00:06:37.920

at a time, because that's not a very, very scalable,

138

00:06:37.920 --> 00:06:41.200

and actually most systems at AWS have kind of long been

139

00:06:41.200 --> 00:06:44.220

beyond the scale that can be managed by hand, you know.

140

00:06:44.220 --> 00:06:48.180

Even the systems at AWS when I joined 13 years ago,

141

00:06:48.180 --> 00:06:51.370

were already quite far beyond that scale.

142

00:06:51.370 --> 00:06:54.410

And we use, you know, today we use the same concepts

143

00:06:54.410 --> 00:06:56.170

you'd be familiar with to manage systems,

144

00:06:56.170 --> 00:06:59.530

we have auto scaling groups, VPC, subnets, security groups,

145

00:06:59.530 --> 00:07:00.870

all of these things.

146

00:07:00.870 --> 00:07:03.550

Internally, we have this concept called host classes,

147

00:07:03.550 --> 00:07:07.750

they're very similar, and they just let us manage, you know,

148

00:07:07.750 --> 00:07:10.710

collections of, of machines without having to think about

149

00:07:10.710 --> 00:07:13.540

the individual machines, and our deployment systems,

150

00:07:13.540 --> 00:07:16.240

you know, they work on top of that, they clone these entire

151

00:07:16.240 --> 00:07:18.403

abstractions between regions, for example.

152

00:07:19.450 --> 00:07:22.730

But we've actually been going much, much further, right?

153

00:07:22.730 --> 00:07:25.790

So, we've learned that we, you know, we have to be able to

154

00:07:25.790 --> 00:07:28.850

operate our systems in hands-off ways,

155

00:07:28.850 --> 00:07:33.850

both for safety, operational safety and for security.

156

00:07:34.120 --> 00:07:37.689

And typically, you know, AWS operators, you know,

157

00:07:37.689 --> 00:07:41.620

like me, as a developer, just simply don't have access

158

00:07:41.620 --> 00:07:44.170

to say, every AWS region.

159

00:07:44.170 --> 00:07:48.500

And so, I could never really plan on being able to log in

160

00:07:48.500 --> 00:07:51.140

or access the system to do something, it's just not our

161

00:07:51.140 --> 00:07:52.663

operational model at all.

162

00:07:54.188 --> 00:07:57.000

But then we have some systems, like AWS outposts

163

00:07:57.000 --> 00:07:59.710

and Snowball that are designed to be disconnected

164

00:07:59.710 --> 00:08:00.700

for periods of time.

165

00:08:00.700 --> 00:08:02.740

You know, we may just have no access to them,

166

00:08:02.740 --> 00:08:05.050

nobody may have any access to them for a while

167

00:08:05.050 --> 00:08:07.730

because of the nature of the product.

168

00:08:07.730 --> 00:08:10.830

And so, we've really had to double and triple down on the,

169

00:08:10.830 --> 00:08:12.900

the automation that it takes to be able to operate

170

00:08:12.900 --> 00:08:15.130

in that environment, to have systems that are self-healing

171

00:08:15.130 --> 00:08:17.593

and that can kind of take care of themselves.

172

00:08:19.160 --> 00:08:22.080

We have kind of two classes of systems these days.

173

00:08:22.080 --> 00:08:26.120

We've kind of got, we've got bastioned systems where we have

174

00:08:26.120 --> 00:08:29.980

very limited access to those systems via what we call

175

00:08:29.980 --> 00:08:32.180

bastion service or bastion hosts

176

00:08:33.500 --> 00:08:37.680

that let us, you know, recover systems typically if,

177

00:08:37.680 --> 00:08:40.750

if there were some urgent need to, or something like that.

178

00:08:40.750 --> 00:08:45.240

But with, you know, a strong record of the fact that folks

179

00:08:45.240 --> 00:08:48.871

have had to either do something with that system and,

180

00:08:48.871 --> 00:08:51.350

and notification processes for that and so on.

181

00:08:51.350 --> 00:08:54.940

But we also have systems where, increasingly that's just

182

00:08:54.940 --> 00:08:55.920

not an option.

183

00:08:55.920 --> 00:08:59.020

You know, there's no general purpose interactive

184

00:08:59.020 --> 00:09:02.900

administrative access, and even in the event of having to

185

00:09:02.900 --> 00:09:05.200

recover a system, there's,

186

00:09:05.200 --> 00:09:08.750

there's just simply no general purpose way to access it.

187

00:09:08.750 --> 00:09:11.680

And instead, we have to have built in in advance, you know,

188

00:09:11.680 --> 00:09:14.650

all the mechanisms that we would need.

189

00:09:14.650 --> 00:09:16.750

You know, the AWS Nitro system is a really strong

190

00:09:16.750 --> 00:09:18.910

example of that, and that's the, that's the system

191

00:09:18.910 --> 00:09:22.450

that runs, you know, our modern EC2 instance types.

192

00:09:22.450 --> 00:09:27.300

And, you know, there's simply no mechanism at all where,

193

00:09:27.300 --> 00:09:30.150

you know, I or someone else could go run a command

194

00:09:30.150 --> 00:09:31.900

on that system or,

195

00:09:31.900 --> 00:09:35.090

or access it in some kind of interactive way.

196

00:09:35.090 --> 00:09:37.640

Instead, everything's done via kind of pre-planned,

197

00:09:38.511 --> 00:09:40.860

pre secured, you know, fully authenticated

198

00:09:40.860 --> 00:09:42.810

and encrypted APIs, and that's the only way

199

00:09:42.810 --> 00:09:44.160

for things to happen.

200

00:09:44.160 --> 00:09:48.140

And, and we found, you know, great motivation for this

201

00:09:48.140 --> 00:09:50.430

is security, but this, this talk is about availability

202

00:09:50.430 --> 00:09:53.280

and a great benefit for availability too, is it just means

203

00:09:53.280 --> 00:09:56.450

there's no possibility of any untracked changes.

204

00:09:56.450 --> 00:09:59.200

And there's no possibility of, you know,

205

00:09:59.200 --> 00:10:01.460

somebody did something to fix something and then forgot

206

00:10:01.460 --> 00:10:02.850

about it or anything like that.

207

00:10:02.850 --> 00:10:04.940

It's just, it can't be that kind of system.

208

00:10:04.940 --> 00:10:07.160

Instead, it's much more hermetic.

209

00:10:07.160 --> 00:10:09.140

So that's our, that's our second item.

210

00:10:09.140 --> 00:10:11.500

I'm gonna hand over to Yasemin, who's gonna now tell us

211

00:10:11.500 --> 00:10:13.843

about the third item on our list.

212

00:10:15.210 --> 00:10:16.347

<v ->Thank you, Colm.</v>

213

00:10:17.650 --> 00:10:21.590

Well, while we design to avoid failures altogether.

214

00:10:21.590 --> 00:10:24.670

we also design with failures in mind.

215

00:10:24.670 --> 00:10:26.540

This is mainly because we,

216

00:10:26.540 --> 00:10:31.060

we know that failures may happen in rare cases,

217

00:10:31.060 --> 00:10:33.940

and you want to make sure that the blast radius

218

00:10:33.940 --> 00:10:38.690

of such events are as minimal as possible.

219

00:10:38.690 --> 00:10:42.670

I will discuss four techniques that we use for this purpose.

220

00:10:42.670 --> 00:10:47.120

The first technique we use is regional isolation.

221

00:10:47.120 --> 00:10:52.100

AWS cloud spans 25 geographic regions around the world,

222

00:10:52.100 --> 00:10:55.020

and there will be eight more regions that are announced

223

00:10:55.020 --> 00:10:56.203

to be launching soon.

224

00:10:59.540 --> 00:11:03.120

Each one of these region is isolated from each other,

225

00:11:03.120 --> 00:11:06.773

both geographically, as well as on the software stack.

226

00:11:07.760 --> 00:11:11.030

AWS services offer regional end points,

227

00:11:11.030 --> 00:11:14.350

enabling direct access to that region.

228

00:11:14.350 --> 00:11:19.330

Using regional isolation, any rare failure that may come up

229

00:11:19.330 --> 00:11:22.940

in one region stays in that region.

230

00:11:22.940 --> 00:11:27.610

Say for example, we have an event going on in us-west-1,

231

00:11:27.610 --> 00:11:29.825

it will be contained within the region,

232

00:11:29.825 --> 00:11:33.523

and will not spread to any of the other regions.

233

00:11:36.310 --> 00:11:40.060

Next technique we use is zonal isolation.

234

00:11:40.060 --> 00:11:43.960

Each AWS region offers multiple availability zones

235

00:11:43.960 --> 00:11:46.217

within it's periphery.

236

00:11:46.217 --> 00:11:50.360

Availability zones are data centers that are miles away

237

00:11:50.360 --> 00:11:51.193

from each other.

238

00:11:52.240 --> 00:11:56.360

AWS cloud has 81 availability zones across the world,

239

00:11:56.360 --> 00:11:59.343

with 24 more launching soon.

240

00:12:00.800 --> 00:12:04.963

Similar to regions, AZ's are also isolated from each other.

241

00:12:05.877 --> 00:12:09.150

A failure in one AZ, let's say us-west-2a

242

00:12:10.020 --> 00:12:12.790

is having an event, will not spread

243

00:12:12.790 --> 00:12:14.043

to any of the other AZ's.

244

00:12:16.740 --> 00:12:19.550

One example where we leverage this property

245

00:12:19.550 --> 00:12:21.550

is regional services.

246

00:12:21.550 --> 00:12:26.200

There will be multiple availability zones serving traffic,

247

00:12:26.200 --> 00:12:28.670

routed to regional end points.

248

00:12:28.670 --> 00:12:32.760

This way, in the event of availability zone failures,

249

00:12:32.760 --> 00:12:36.570

traffic will be redirected to healthy zones

250

00:12:36.570 --> 00:12:40.690

behind the region, and customers will not be serving,

251

00:12:40.690 --> 00:12:43.560

will not be experiencing any degradation.

252

00:12:43.560 --> 00:12:47.290

So far, we discussed physical compartmentalization

253

00:12:47.290 --> 00:12:52.290

of architecture by using AWS regions and availability zones.

254

00:12:52.410 --> 00:12:55.480

These two have regional and zonal blast radius

255

00:12:55.480 --> 00:12:56.713

impact accordingly.

256

00:12:58.640 --> 00:13:01.510

We apply cellular isolation to further reduce

257

00:13:01.510 --> 00:13:04.020

the blast radius impact of events.

258

00:13:04.020 --> 00:13:07.540

In this technique, we build dedicated software stacks

259

00:13:07.540 --> 00:13:09.320

which we call cells.

260

00:13:09.320 --> 00:13:11.860

Cells are isolated from each other,

261

00:13:11.860 --> 00:13:14.103

with their own dedicated end points.

262

00:13:15.370 --> 00:13:18.270

Customers get assigned to one of these cells.

263

00:13:18.270 --> 00:13:21.300

Imagine there are eight customers and four cells,

264

00:13:21.300 --> 00:13:23.693

we have two customers assigned to each of them.

265

00:13:24.840 --> 00:13:27.370

Since cells are isolated from each other,

266

00:13:27.370 --> 00:13:31.210

failure in one cell, let's say cell two is having an event,

267

00:13:31.210 --> 00:13:33.870

will not spread to any of the other cells.

268

00:13:33.870 --> 00:13:36.750

That means it's only the customers assigned to that cell,

269

00:13:36.750 --> 00:13:40.350

the palm tree and stars in this case, will be impacted,

270

00:13:40.350 --> 00:13:43.063

but no other customers will be impacted.

271

00:13:44.240 --> 00:13:45.890

This is pretty good, right?

272

00:13:45.890 --> 00:13:50.000

It's much better than a zonal and regional blast radius

273

00:13:50.000 --> 00:13:51.940

that we discussed before.

274

00:13:51.940 --> 00:13:53.093

But can we do better?

275

00:13:54.280 --> 00:13:58.050

We do better by using shuffle sharding.

276

00:13:58.050 --> 00:14:01.730

In shuffle sharding, we divide the service into smaller

277

00:14:01.730 --> 00:14:05.230

compartments, which I will call partitions in this case,

278

00:14:05.230 --> 00:14:06.473

the blue boxes.

279

00:14:07.624 --> 00:14:11.680

And we assign customers to partitions, not one to one,

280

00:14:11.680 --> 00:14:14.630

but we assign multiple, multiple partitions

281

00:14:14.630 --> 00:14:16.270

to a single customer.

282

00:14:16.270 --> 00:14:19.700

In this case, I have two partitions assigned

283

00:14:19.700 --> 00:14:21.760

to each customer.

284

00:14:21.760 --> 00:14:25.130

So, let's evaluate the blast radius impact of failures

285

00:14:25.130 --> 00:14:26.650

in this case.

286

00:14:26.650 --> 00:14:30.690

Let's say partition one is having an event.

287

00:14:30.690 --> 00:14:34.620

In that case, both palm tree and stars have another

288

00:14:34.620 --> 00:14:36.540
partition that's healthy.

289

00:14:36.540 --> 00:14:41.030
Therefore they will not see impact and they will continue to

290

00:14:41.030 --> 00:14:42.293
operate just fine.

291

00:14:44.220 --> 00:14:47.980
It's only the time around both partitions of a single

292

00:14:47.980 --> 00:14:51.490
customer having an event then that customer

293

00:14:51.490 --> 00:14:53.350
will see the impact.

294

00:14:53.350 --> 00:14:56.980
And in this case, when partition four is also out,

295

00:14:56.980 --> 00:15:01.280
it's the stars that will see the impact because those two

296

00:15:01.280 --> 00:15:03.530
are shared partition for stars,

297

00:15:03.530 --> 00:15:07.793
but palm tree and hiker will not observe any impact.

298

00:15:09.600 --> 00:15:13.580
So there are two benefits that I like to highlight here.

299

00:15:13.580 --> 00:15:17.040
The real first one is creating customer impact

300

00:15:17.040 --> 00:15:19.670

is much harder with shuffle sharding.

301

00:15:19.670 --> 00:15:23.410

This is mainly because it takes multiple partitions

302

00:15:23.410 --> 00:15:26.253

to have failure, to create an impact,

303

00:15:27.110 --> 00:15:30.560

and probability of having a failure across multiple

304

00:15:30.560 --> 00:15:34.900

partitions is much lower than probability of having

305

00:15:34.900 --> 00:15:37.113

a failure on a single partition.

306

00:15:38.410 --> 00:15:42.920

The second benefit is that even in the amount of those very,

307

00:15:42.920 --> 00:15:46.110

very rare probability events that are happening

308

00:15:46.110 --> 00:15:50.396

across partitions, the impact created to the customer

309

00:15:50.396 --> 00:15:52.410

is much lower.

310

00:15:52.410 --> 00:15:55.630

We have one customer being impacted in the case of

311

00:15:55.630 --> 00:15:59.640

two partition failures compared to impacting multiple

312

00:15:59.640 --> 00:16:02.650

customers in the regular sharding schemes

313

00:16:02.650 --> 00:16:03.990
would actually create.

314

00:16:03.990 --> 00:16:06.820
Next, I will discuss circuit breakers.

315

00:16:06.820 --> 00:16:09.740
Circuit breakers is one of those techniques we used to

316

00:16:09.740 --> 00:16:13.373
eliminate failures, and I will discuss two examples.

317

00:16:14.410 --> 00:16:15.883
First one is load shedding.

318

00:16:16.720 --> 00:16:21.610
We know systems can slow down, or sometimes even fall over

319

00:16:21.610 --> 00:16:23.530
under excess load.

320

00:16:23.530 --> 00:16:26.320
We design to make sure that our services are not

321

00:16:26.320 --> 00:16:27.773
vulnerable to this problem.

322

00:16:28.860 --> 00:16:32.720
To address this issue, we first identify the maximum

323

00:16:32.720 --> 00:16:35.683
capacity of every individual component.

324

00:16:36.770 --> 00:16:39.980
We use stress testing to get this information.

325

00:16:39.980 --> 00:16:43.610

Now we install the load shedders locally within the services

326

00:16:43.610 --> 00:16:45.903

to monitor the traffic being served.

327

00:16:47.130 --> 00:16:49.130

Once service starts receiving traffic

328

00:16:49.130 --> 00:16:52.560

more than it's predefined maximum capacity limit,

329

00:16:52.560 --> 00:16:56.293

the load shedders start rejecting any excess load.

330

00:16:57.260 --> 00:16:59.920

They are designed to reject a load very quickly

331

00:16:59.920 --> 00:17:03.160

without spending much of the system's resources.

332

00:17:03.160 --> 00:17:06.900

This way we ensure that systems continue to operate

333

00:17:06.900 --> 00:17:09.193

successfully under excess load.

334

00:17:12.490 --> 00:17:14.930

The other circuit breaker that I'll discuss

335

00:17:14.930 --> 00:17:16.880

is bullet counters.

336

00:17:16.880 --> 00:17:20.460

This pattern gets used frequently on external monitoring

337

00:17:20.460 --> 00:17:24.160

applications, like Heart Guardians in this case.

338

00:17:24.160 --> 00:17:28.730

These applications are common across AWS because we want to

339

00:17:28.730 --> 00:17:32.910

make sure that we are the first ones to detect any problem

340

00:17:32.910 --> 00:17:36.520

that's maybe ongoing and mitigate it right away

341

00:17:36.520 --> 00:17:38.483

before customers are even noticing it.

342

00:17:40.150 --> 00:17:44.410

For example, these applications pair the cloud checks

343

00:17:44.410 --> 00:17:48.040

to each and every node of the system to ensure that

344

00:17:48.040 --> 00:17:49.690

they're all healthy.

345

00:17:49.690 --> 00:17:52.190

If there's a bad node being detected,

346

00:17:52.190 --> 00:17:54.270

they would replace the node.

347

00:17:54.270 --> 00:17:57.150

Well, it sounds straightforward, right?

348

00:17:57.150 --> 00:18:00.430

Well the part that's interesting about these systems

349

00:18:00.430 --> 00:18:03.880

is that they are very, very powerful.

350

00:18:03.880 --> 00:18:07.480

Replacing a node is a powerful action.

351

00:18:07.480 --> 00:18:11.070

Imagine that Health Guardian determined that half of the

352

00:18:11.070 --> 00:18:13.290

fleet is unhealthy.

353

00:18:13.290 --> 00:18:17.140

Should it go ahead and just replace all those nodes?

354

00:18:17.140 --> 00:18:19.880

Maybe, but usually not.

355

00:18:19.880 --> 00:18:22.723

It really depends on what the problem is.

356

00:18:23.840 --> 00:18:28.360

We built to not have such large failures anyways, right?

357

00:18:28.360 --> 00:18:33.357

So to avoid these automations to take

358

00:18:34.950 --> 00:18:38.973

significant actions that will lead to significant changes,

359

00:18:40.060 --> 00:18:44.020

we installed the bullet counters in place

360

00:18:44.020 --> 00:18:47.963

so that they don't act on (indistinct) signals incorrectly.

361

00:18:49.500 --> 00:18:53.810

So bullet counter in this example would be the one that we

362

00:18:53.810 --> 00:18:56.970

say, "What's the maximum percentage of the fleet

363

00:18:56.970 --> 00:19:00.287

that could be replaced safely at a given time?"

364

00:19:02.120 --> 00:19:06.250

The bullet counters monitor the actions taken

365

00:19:06.250 --> 00:19:11.250

by the Health Guardian, and whenever there is an event where

366

00:19:12.457 --> 00:19:15.200

there are nodes that are being determined that's more than

367

00:19:15.200 --> 00:19:19.960

predefined maximum limit, they will stop the execution,

368

00:19:19.960 --> 00:19:24.520

and instead they will notify the operators to show up

369

00:19:24.520 --> 00:19:26.143

and assess the situation.

370

00:19:28.060 --> 00:19:31.330

Using bullet counters, we ensure that automation

371

00:19:31.330 --> 00:19:34.733

always operates within known and safe limits.

372

00:19:38.040 --> 00:19:40.283

With that, I will hand it over to Colm.

373

00:19:41.480 --> 00:19:42.543

<v ->Thanks, Yasemin.</v>

374

00:19:43.790 --> 00:19:46.180

So as I mentioned briefly earlier,

375

00:19:46.180 --> 00:19:49.090

testing is of enormous importance,

376

00:19:49.090 --> 00:19:52.620

and we invest a lot in testing.

377

00:19:52.620 --> 00:19:56.460

In fact, for our, our most highly available systems,

378

00:19:56.460 --> 00:19:59.360

it's not unusual to spend, you know,

379

00:19:59.360 --> 00:20:04.360

much more time on testing code than writing the code itself.

380

00:20:05.000 --> 00:20:09.390

You know, it's, we have unit tests,

381

00:20:09.390 --> 00:20:12.700

we've got end-to-end tests, we've got integration tests,

382

00:20:12.700 --> 00:20:14.773

and we've even got formal verification.

383

00:20:15.690 --> 00:20:19.643

It's, we've got, you know, pre-production environments.

384

00:20:20.570 --> 00:20:25.470

We, even these days, you know, we'll test how a system

385

00:20:25.470 --> 00:20:28.540

cooperates with rolling forward to a deployment process,

386

00:20:28.540 --> 00:20:30.950

as well as rolling back through a deployment process

387

00:20:30.950 --> 00:20:33.150

before we ever really deploy to production,

388

00:20:33.150 --> 00:20:34.630
just to make sure that like, well,

389

00:20:34.630 --> 00:20:38.083
if we had to do those things, we would be able to.

390

00:20:39.110 --> 00:20:42.234
If you want to see examples of,

391

00:20:42.234 --> 00:20:44.253
of how we raise the bar on testing,

392

00:20:45.188 --> 00:20:47.970
you can look at our open source projects on GitHub.

393

00:20:47.970 --> 00:20:52.970
So our s2n project, which is our open source SSL and TLS

394

00:20:53.220 --> 00:20:58.220
library is developed by a team at Amazon who work just like

395

00:20:58.480 --> 00:21:01.160
any other team in Amazon, and it's a really simple,

396

00:21:01.160 --> 00:21:05.068
like open view into, into how we work.

397

00:21:05.068 --> 00:21:07.748
And you can see just the sheer staggering number of tests

398

00:21:07.748 --> 00:21:12.100
that are there, that are running on every single build.

399

00:21:12.100 --> 00:21:14.530
You know, when we don't, we don't just run those tests

400

00:21:14.530 --> 00:21:19.530

before we deploy, they're integrated into our CICD process,

401

00:21:20.140 --> 00:21:23.940

and, you know, every time we check in, we run them all,

402

00:21:23.940 --> 00:21:26.323

and we make sure that they all pass before,

403

00:21:27.180 --> 00:21:28.720

before promoting, just to make sure

404

00:21:28.720 --> 00:21:29.880

that there are no regressions.

405

00:21:29.880 --> 00:21:32.140

And that's very, very typical for these

406

00:21:32.140 --> 00:21:34.390

high availability systems.

407

00:21:34.390 --> 00:21:37.417

But, we've been going further.

408

00:21:37.417 --> 00:21:40.860

An so, in the last few years, we've actually been going

409

00:21:40.860 --> 00:21:43.070

further and further with automated reasoning

410

00:21:43.070 --> 00:21:45.060

and formal verification.

411

00:21:45.060 --> 00:21:48.060

So this is something we've been doing for a long time,

412

00:21:48.060 --> 00:21:51.840

you know, on the s2n project, we've been formally verifying

413

00:21:51.840 --> 00:21:56.833

the correctness of parts of s2n for about six years now.

414

00:21:57.800 --> 00:22:00.450

But in, in the last few years,

415

00:22:00.450 --> 00:22:03.293

we've been able to get to the point where now, you know,

416

00:22:05.660 --> 00:22:07.950

I'm gonna say regular developers who don't have

417

00:22:07.950 --> 00:22:10.320

specialized training in automated reasoning

418

00:22:10.320 --> 00:22:13.870

or formal verification are able to use these techniques

419

00:22:13.870 --> 00:22:15.690

because we've improved the tooling,

420

00:22:15.690 --> 00:22:17.343

it's getting easier and easier.

421

00:22:18.350 --> 00:22:20.760

And this is amazing because formal verification, you know,

422

00:22:20.760 --> 00:22:24.670

tests can always prove that code is correct for a particular

423

00:22:24.670 --> 00:22:27.440

input, but formal verification and can prove that that code

424

00:22:27.440 --> 00:22:29.290

is correct for any particular input.

425

00:22:29.290 --> 00:22:31.310

So it's very, very powerful.

426

00:22:31.310 --> 00:22:35.090

We can find, you know, really hard to find edge cases

427

00:22:35.090 --> 00:22:37.490

that you won't find using any other technique.

428

00:22:37.490 --> 00:22:39.310

And it's pretty awesome now that, you know,

429

00:22:39.310 --> 00:22:43.290

we're able to get to the point where, you know,

430

00:22:43.290 --> 00:22:45.540

regular software developers like me can,

431

00:22:45.540 --> 00:22:49.980

can actually use these tools and improve things

432

00:22:49.980 --> 00:22:50.813

about their code.

433

00:22:50.813 --> 00:22:51.900

And if you're interested in some of that,

434

00:22:51.900 --> 00:22:54.340

I'd encourage you to check out CBMC,

435

00:22:54.340 --> 00:22:58.310

which is a tool we've been using on s2n that is very

436

00:22:58.310 --> 00:23:03.310

intuitive, and I found very easy to use for developers.

437

00:23:03.400 --> 00:23:08.020

So, the next item on our list is something called

438

00:23:08.020 --> 00:23:12.720

lifecycle management, which we found that, you know,

439

00:23:12.720 --> 00:23:16.350

to get high availability over long periods of time

440

00:23:16.350 --> 00:23:19.820

in our systems, we need to be very, very intentional about

441

00:23:19.820 --> 00:23:23.620

how we manage the life cycle of a lot of aspects

442

00:23:23.620 --> 00:23:24.623

of those systems.

443

00:23:25.853 --> 00:23:29.730

And in particular, any kind of credentials that are used

444

00:23:29.730 --> 00:23:30.563

by those systems.

445

00:23:30.563 --> 00:23:33.920

So that's things like keys and certificates and so on.

446

00:23:33.920 --> 00:23:36.700

So modern security and compliance frameworks demand that

447

00:23:36.700 --> 00:23:39.000

credentials be frequently rotated, and,

448

00:23:39.000 --> 00:23:41.780

and that makes sense, no one wants a key or a certificate

449

00:23:41.780 --> 00:23:44.580

around that could be used for very long periods of time.

450

00:23:45.590 --> 00:23:48.700

But at the same time, expired and mismatched credentials

451

00:23:48.700 --> 00:23:50.430

could be a source of outages, right?

452

00:23:50.430 --> 00:23:52.770

If a certificate or key expires and it's still

453

00:23:52.770 --> 00:23:55.260

in your system, that's not gonna be good.

454

00:23:55.260 --> 00:23:58.920

So we've learned to decouple that expiry from alarming.

455

00:23:58.920 --> 00:24:02.930

So we, we alarm well before the expiry time of any key

456

00:24:02.930 --> 00:24:07.000

or credential, and we've learned to be super intentional

457

00:24:07.000 --> 00:24:11.800

and kind of go into overkill on how to monitor it.

458

00:24:11.800 --> 00:24:14.530

So we've got time to expiry metrics for anything

459

00:24:14.530 --> 00:24:15.800

that expires.

460

00:24:15.800 --> 00:24:18.700

We look at that from both the server side perspective,

461

00:24:18.700 --> 00:24:20.620

so the thing that might be serving a certificate

462

00:24:20.620 --> 00:24:23.090

or using a key, and the client side,

463

00:24:23.090 --> 00:24:25.083
thing that's connecting to it.

464

00:24:25.083 --> 00:24:27.540
Like I said, we alarm and investigate well before

465

00:24:27.540 --> 00:24:29.220
there's any kind of a problem.

466

00:24:29.220 --> 00:24:32.500
And, and then on top of that, we've got additional

467

00:24:32.500 --> 00:24:35.340
fail safes, and canaries that are constantly scanning

468

00:24:35.340 --> 00:24:38.370
for anything that looks like it's even close to expiry,

469

00:24:38.370 --> 00:24:41.080
so that we've got another safety net there too.

470

00:24:41.080 --> 00:24:43.760
And then, with all of those systems combined,

471

00:24:43.760 --> 00:24:45.090
we've learned that, you know,

472

00:24:45.090 --> 00:24:48.650
that means we need to deploy any new key or credential,

473

00:24:48.650 --> 00:24:51.030
make sure it's absolutely everywhere that could need it

474

00:24:51.030 --> 00:24:53.787
before we activate it, and only then activate it

475

00:24:53.787 --> 00:24:55.820

and kind of do the same in reverse whenever

476

00:24:55.820 --> 00:24:59.290

we're deactivating or evoking a key or certificate.

477

00:24:59.290 --> 00:25:00.830

And we found that, you know,

478

00:25:00.830 --> 00:25:03.810

paying particular attention to detail on those processes

479

00:25:03.810 --> 00:25:07.660

has been really key to avoiding just any kind of outage

480

00:25:07.660 --> 00:25:11.630

you could see from just expired or mismatched keys

481

00:25:11.630 --> 00:25:13.200

and credentials.

482

00:25:13.200 --> 00:25:16.210

And so with that, I'll hand over to Yasemin,

483

00:25:16.210 --> 00:25:19.050

who's gonna tell us about modular separations.

484

00:25:19.050 --> 00:25:20.207

<v ->Thank you, Colm.</v>

485

00:25:22.036 --> 00:25:23.253

Modular separation.

486

00:25:24.270 --> 00:25:27.300

We avoid multi-link architectures and decouple

487

00:25:27.300 --> 00:25:29.724

individual responsibilities into their own

488

00:25:29.724 --> 00:25:31.474

dedicated components.

489

00:25:33.480 --> 00:25:35.870

Control plane versus Data plane separation

490

00:25:35.870 --> 00:25:37.800

is a good example of this.

491

00:25:37.800 --> 00:25:42.027

Most data services have this notion of control plane

492

00:25:42.027 --> 00:25:44.100

and data plane APIs.

493

00:25:44.100 --> 00:25:46.733

Let's take Kinesis data streams as an example.

494

00:25:47.800 --> 00:25:51.270

Kinesis data streams is a real time streaming service

495

00:25:51.270 --> 00:25:54.780

that enables customers to write records into a log stream,

496

00:25:54.780 --> 00:25:55.973

and read them later on.

497

00:25:57.044 --> 00:25:58.820

And the very first thing.

498

00:25:58.820 --> 00:26:02.070

The very first thing customers do when they start using

499

00:26:02.070 --> 00:26:05.520

the service, they go ahead and create a stream.

500

00:26:05.520 --> 00:26:08.147

Create stream is a control plane API.

501

00:26:09.003 --> 00:26:10.700

Once the stream is created,

502

00:26:10.700 --> 00:26:13.470

producer applications can start ingesting the data

503

00:26:13.470 --> 00:26:17.300

continuously and consumer applications read those records

504

00:26:17.300 --> 00:26:18.453

within milliseconds.

505

00:26:19.960 --> 00:26:23.716

The APIs used by the producer and consumer applications

506

00:26:23.716 --> 00:26:25.293

are the data plane APIs.

507

00:26:26.396 --> 00:26:27.696

As we see in this example,

508

00:26:28.687 --> 00:26:33.687

access patterns of these two, these two type of APIs

509

00:26:34.270 --> 00:26:35.103

are different.

510

00:26:36.470 --> 00:26:40.500

Consequently, their dependencies are very different as well.

511

00:26:40.500 --> 00:26:44.900

For example, control plane APIs are depending on the

512

00:26:44.900 --> 00:26:49.900

asynchronous workflows to execute the steps of creating

513

00:26:50.390 --> 00:26:55.390

a stream, versus data plane APIs are dependent on the data

514

00:26:55.804 --> 00:26:58.433

store to fetch the records within milliseconds.

515

00:26:59.970 --> 00:27:04.000

By decoupling the two API types, we limit the blast,

516

00:27:04.000 --> 00:27:07.970

we limit the impact their dependencies could be creating.

517

00:27:07.970 --> 00:27:10.950

For example, if there's an event that's going on

518

00:27:10.950 --> 00:27:14.170

in the Async workflows, then it's going to be only

519

00:27:14.170 --> 00:27:16.850

the control plane API's that are being impacted

520

00:27:16.850 --> 00:27:19.713

while data plane APIs continue to work.

521

00:27:23.550 --> 00:27:26.260

Next technique is static stability.

522

00:27:26.260 --> 00:27:28.540

And this technique works hand to hand

523

00:27:28.540 --> 00:27:29.740

with modular separation.

524

00:27:32.020 --> 00:27:35.830

Availability of a system can be as good as

525

00:27:35.830 --> 00:27:37.853

availability of its dependencies.

526

00:27:38.830 --> 00:27:41.820

For this reason, we strive to keep the dependencies of

527

00:27:41.820 --> 00:27:43.663

systems to absolute minimum.

528

00:27:46.520 --> 00:27:49.400

Let's get back to control plane versus data plane

529

00:27:49.400 --> 00:27:50.363

discussion again.

530

00:27:51.710 --> 00:27:54.400

The dependency here that I didn't discuss before

531

00:27:54.400 --> 00:27:56.190

is the Metadata store.

532

00:27:56.190 --> 00:27:59.989

Metadata store is the one that process the information

533

00:27:59.989 --> 00:28:03.023

about resources, Kinesis streams in this case.

534

00:28:04.070 --> 00:28:08.200

So stream creation execution has direct dependency

535

00:28:08.200 --> 00:28:11.240

to the availability of Metadata store because it needs

536

00:28:11.240 --> 00:28:14.700

to process the information that the stream is created.

537

00:28:14.700 --> 00:28:18.463

Imagine there's an outage on Metadata store.

538

00:28:18.463 --> 00:28:20.750

We expect the Async workflows to be impacted

539

00:28:20.750 --> 00:28:23.230

because it cannot access the data store anymore,

540

00:28:23.230 --> 00:28:26.280

therefore the control plane APIs are impacted.

541

00:28:26.280 --> 00:28:30.500

Let's check the impact of Metadata store on data plane APIs.

542

00:28:31.720 --> 00:28:36.720

Well, data plane APIs also need to know about these streams,

543

00:28:36.977 --> 00:28:39.590

and where they are located so that they can serve,

544

00:28:39.590 --> 00:28:41.383

put and get APIs on them.

545

00:28:42.360 --> 00:28:45.010

So if the data plane APIs have direct and synchronous

546

00:28:45.010 --> 00:28:49.690

dependency on Metadata store, then its outage

547

00:28:49.690 --> 00:28:51.603

will also impact the data plane APIs.

548

00:28:53.300 --> 00:28:54.910

Well, that's not great.

549

00:28:54.910 --> 00:28:58.120

It's large blast radius impact.

550

00:28:58.120 --> 00:29:01.250

So, the sole dependency of data plane APIs

551

00:29:01.250 --> 00:29:03.650

is really the data store, they are responsible

552

00:29:03.650 --> 00:29:06.660

for serving customer records.

553

00:29:06.660 --> 00:29:10.610

So, considering the static stability principle,

554

00:29:10.610 --> 00:29:14.970

let's look at this design and see how we can eliminate

555

00:29:14.970 --> 00:29:18.623

the Metadata store dependency from data plane APIs.

556

00:29:19.940 --> 00:29:24.923

We eliminate this dependency by moving the Metadata store

557

00:29:24.923 --> 00:29:29.818

dependency of data plane from being a synchronous dependency

558

00:29:29.818 --> 00:29:31.793

to being an asynchronous dependency.

559

00:29:33.010 --> 00:29:38.010

In this architecture, the Metadata store keeps track

560

00:29:38.170 --> 00:29:40.293

of the stream of information.

561

00:29:41.190 --> 00:29:46.040

And here it has a copy of this information

562

00:29:46.040 --> 00:29:48.183

stored within the data plane itself.

563

00:29:51.420 --> 00:29:54.667

Any updates that are being applied to this metadata

564

00:29:54.667 --> 00:29:56.920

gets propagated to the data plane,

565

00:29:56.920 --> 00:29:59.653

and it's own snapshot asynchronously.

566

00:30:01.190 --> 00:30:04.770

In this case, if there's an outage in Metadata store,

567

00:30:04.770 --> 00:30:08.240

it's the control plane APIs those will be impacted,

568

00:30:08.240 --> 00:30:11.520

but the data plane APIs will continue to work

569

00:30:11.520 --> 00:30:15.790

by using the smaller version of their snapshots

570

00:30:15.790 --> 00:30:17.893

that are stored within themselves.

571

00:30:19.660 --> 00:30:23.980

This way, we ensure that the data plane is statically stable

572

00:30:23.980 --> 00:30:27.733

by using its absolute minimum set of dependencies.

573

00:30:28.730 --> 00:30:29.760

<v ->Thanks Yasemin.</v>

574

00:30:31.390 --> 00:30:35.640

So our ninth and penultimate item is constant work,

575

00:30:35.640 --> 00:30:36.473

which is,

576

00:30:37.920 --> 00:30:40.210

you know, whenever we were learning big O notation,

577

00:30:40.210 --> 00:30:42.410

if you study computer science, you know,

578

00:30:42.410 --> 00:30:44.610

you probably learnt about systems that are constant of work,

579

00:30:44.610 --> 00:30:46.670

that means O 1, right?

580

00:30:46.670 --> 00:30:48.570

A big O notation.

581

00:30:48.570 --> 00:30:50.710

And that turns out to be a very important concept

582

00:30:50.710 --> 00:30:53.860

for some of our highly available systems.

583

00:30:53.860 --> 00:30:56.990

You know, in general, right?

584

00:30:56.990 --> 00:31:00.970

Risk is proportionate to rates of change in systems, right?

585

00:31:00.970 --> 00:31:03.680

A spike in load, for example, right?

586

00:31:03.680 --> 00:31:05.710

Can cause the system to slow down, right?

587

00:31:05.710 --> 00:31:09.070

And then the system gets into a mode that it's not used to

588

00:31:09.070 --> 00:31:11.560

operating in, things don't really know how to handle that,

589

00:31:11.560 --> 00:31:13.630

they might start timing out and so on.

590

00:31:13.630 --> 00:31:16.550

And the, you know, issues like that can cause

591

00:31:16.550 --> 00:31:18.380

cascading failures.

592

00:31:18.380 --> 00:31:21.890

And so we've learned that reducing the overall dynamism

593

00:31:21.890 --> 00:31:24.090

in the system, you know, the amount of change it can,

594

00:31:24.090 --> 00:31:25.530

it can even go through,

595

00:31:25.530 --> 00:31:28.110

is a useful way to make them simpler and to reduce

596

00:31:28.110 --> 00:31:29.230

all that risk.

597

00:31:29.230 --> 00:31:31.810

And a kind of counter-intuitive solution to kind of wrangle

598

00:31:31.810 --> 00:31:35.060

that risk, is actually to run the system at maximum load

599

00:31:35.060 --> 00:31:36.410

all the time, right?

600

00:31:36.410 --> 00:31:38.170

And even though that sounds like, "Well, now the system's

601

00:31:38.170 --> 00:31:40.250
gonna be maxed out all the time."

602

00:31:40.250 --> 00:31:42.850
It actually reduces the amount of dynamism and change

603

00:31:42.850 --> 00:31:45.200
that's in the system, and therefore risk.

604

00:31:45.200 --> 00:31:48.210
A really simple example of that is, how we apply

605

00:31:48.210 --> 00:31:53.000
this constant work pattern to say configuration changes.

606

00:31:53.000 --> 00:31:55.940
So a really common pattern and how developers manage

607

00:31:57.070 --> 00:31:58.800
configuration changes is like, you know,

608

00:31:58.800 --> 00:32:01.770
customer makes a change, and that change gets ingested

609

00:32:01.770 --> 00:32:04.870
into the system as like a Delta, you know, do this thing,

610

00:32:04.870 --> 00:32:07.777
do X, do Y, and that goes into a workflow,

611

00:32:07.777 --> 00:32:10.730
and the workflow manages getting that change out to all of

612

00:32:10.730 --> 00:32:13.320
the systems that need to reflect that change.

613

00:32:13.320 --> 00:32:16.660

That works, that's a, a simple pattern,

614

00:32:16.660 --> 00:32:18.670

but the problem is, when lots of changes happen,

615

00:32:18.670 --> 00:32:20.480

so lots of customers at the same time,

616

00:32:20.480 --> 00:32:22.450

maybe it's a particularly busy day or whatever,

617

00:32:22.450 --> 00:32:24.220

the overall system slows down, right?

618

00:32:24.220 --> 00:32:27.730

Because the workflow's got more work to do, right?

619

00:32:27.730 --> 00:32:31.420

And as I said, it could, that can, you know, hit you in,

620

00:32:31.420 --> 00:32:33.420

in ways that cascade.

621

00:32:33.420 --> 00:32:35.430

So a simpler version of this is imagine,

622

00:32:35.430 --> 00:32:37.680

well every customer, they make their change,

623

00:32:37.680 --> 00:32:42.140

and it's just reflected as say, a file or a key in S3,

624

00:32:42.140 --> 00:32:43.140

that's it, right?

625

00:32:43.140 --> 00:32:46.470

They just make their change and it's effectively a file

626

00:32:46.470 --> 00:32:48.460

or a key in S3.

627

00:32:48.460 --> 00:32:51.310

And the system on the other side, instead of using

628

00:32:51.310 --> 00:32:54.470

a workflow, all it's doing is checking all of those files

629

00:32:54.470 --> 00:32:55.490

every single time.

630

00:32:55.490 --> 00:32:58.047

Just downloading every single file from S3

631

00:32:58.047 --> 00:33:00.670

and using that as its configuration.

632

00:33:00.670 --> 00:33:03.190

And so, even when lots and lots of customers make changes

633

00:33:03.190 --> 00:33:06.530

at the same time, so maybe a hundred files changed.

634

00:33:06.530 --> 00:33:08.900

If the system was always pulling, you know,

635

00:33:08.900 --> 00:33:11.600

all 100 files or all thousand files,

636

00:33:11.600 --> 00:33:13.580

if there's a thousand customers, and just doing that

637

00:33:13.580 --> 00:33:17.010

as its configuration, there's no change, or difference,

638

00:33:17.010 --> 00:33:20.800

or dynamism on the right-hand side of these diagrams,

639

00:33:20.800 --> 00:33:22.530

which reduces the overall risk,

640

00:33:22.530 --> 00:33:24.610

and it could be enormously effective.

641

00:33:24.610 --> 00:33:26.020

And we've learned to apply this pattern

642

00:33:26.020 --> 00:33:28.220

in some really key places.

643

00:33:28.220 --> 00:33:31.170

We've got a builder's library article where we talk about

644

00:33:31.170 --> 00:33:33.260

how we apply this in our health check systems,

645

00:33:33.260 --> 00:33:36.000

in our DNS fail-over systems, so that they can be

646

00:33:36.000 --> 00:33:38.210

incredibly reliable.

647

00:33:38.210 --> 00:33:40.730

And I found it really, really useful.

648

00:33:40.730 --> 00:33:44.560

And with that, Yasemin's gonna take over and close us out

649

00:33:44.560 --> 00:33:47.623

and, and tell us our final lesson.

650

00:33:48.751 --> 00:33:49.950

<v ->Retries.</v>

651

00:33:49.950 --> 00:33:51.830

Retries are somewhat well-known, right?

652

00:33:51.830 --> 00:33:56.210

When there is a failure, we retry and that helps

653

00:33:56.210 --> 00:33:57.950

resolving the problem.

654

00:33:57.950 --> 00:33:59.453

But do they always help?

655

00:34:00.300 --> 00:34:02.420

Let's dig in a little bit.

656

00:34:02.420 --> 00:34:05.370

The problem with retries is that when they are not

657

00:34:05.370 --> 00:34:09.060

used properly, they can cause a larger event

658

00:34:09.060 --> 00:34:12.110

compared to what they're trying to mitigate.

659

00:34:12.110 --> 00:34:14.820

We call this thundering herd problem.

660

00:34:14.820 --> 00:34:17.071

I'll first explain what this problem is,

661

00:34:17.071 --> 00:34:21.870

and then I'll discuss two techniques to avoid it.

662

00:34:21.870 --> 00:34:24.160

Imagine there are transient failures going on

663

00:34:24.160 --> 00:34:25.053
in the system.

664

00:34:26.050 --> 00:34:28.320
Clients start to retry.

665

00:34:28.320 --> 00:34:31.470
As clients retry more, there's more traffic being

666

00:34:31.470 --> 00:34:35.150
generated, so system gets overloaded, and when systems get

667

00:34:35.150 --> 00:34:38.730
overloaded, we discussed, they start load shedding,

668

00:34:38.730 --> 00:34:41.310
and there are more failures, and there are more retries,

669

00:34:41.310 --> 00:34:45.093
and this is, it's this vicious cycle that's keep going.

670

00:34:46.120 --> 00:34:48.940
Creating more work in the system with retries

671

00:34:48.940 --> 00:34:51.570
doesn't really have to solve the problem.

672

00:34:51.570 --> 00:34:54.203
So how do we avoid the thundering herd issue?

673

00:34:55.170 --> 00:34:57.270
The first technique I'll discuss

674

00:34:57.270 --> 00:35:00.640
is exponential backoff and jitter.

675

00:35:00.640 --> 00:35:04.650

When clients retry without exponential backoff and jitter,

676

00:35:04.650 --> 00:35:08.640

the same amount of traffic hits the service right around

677

00:35:08.640 --> 00:35:12.610

the same time with the same frequency.

678

00:35:12.610 --> 00:35:17.480

In this example, there's a second delay between each retry,

679

00:35:17.480 --> 00:35:21.403

and the same set of calls are being hit.

680

00:35:22.260 --> 00:35:24.493

Calls are used to hit the service.

681

00:35:25.830 --> 00:35:28.380

So I have two clients represented here

682

00:35:28.380 --> 00:35:32.530

with the shades of colors, and first line is doing one TPS,

683

00:35:32.530 --> 00:35:35.860

second one is doing three, and the third one is doing five

684

00:35:35.860 --> 00:35:38.113

TPS transactions per second, right?

685

00:35:39.750 --> 00:35:42.423

They're hitting service again and again.

686

00:35:43.660 --> 00:35:48.660

So, this doesn't help to lower the load on the service.

687

00:35:49.860 --> 00:35:53.146

With exponential backoff and jitter being used,

688

00:35:53.146 --> 00:35:54.750

we apply two techniques.

689

00:35:54.750 --> 00:35:59.750

The first one is, we give wait time between each retry

690

00:36:01.380 --> 00:36:04.450

and that wait time increases exponentially

691

00:36:04.450 --> 00:36:06.400

in between every retries.

692

00:36:06.400 --> 00:36:09.090

And the second aspect of this is jitter.

693

00:36:09.090 --> 00:36:12.820

We add randomness so that when the retry comes on,

694

00:36:12.820 --> 00:36:17.820

it will come in on that, the retries will not come in

695

00:36:18.130 --> 00:36:21.403

right at the same second, but rather it will be distributed

696

00:36:21.403 --> 00:36:23.023

across the timeframe.

697

00:36:24.630 --> 00:36:26.330

This technique helps,

698

00:36:26.330 --> 00:36:29.223

but we find it's not always sufficient at scale.

699

00:36:33.090 --> 00:36:35.210

Let's look at client throttling.

700

00:36:35.210 --> 00:36:37.130

We found client throttling to be much more

701

00:36:37.130 --> 00:36:38.423

effective technique.

702

00:36:39.950 --> 00:36:43.150

AWS the case have built in support for client throttling

703

00:36:43.150 --> 00:36:44.470

as well.

704

00:36:44.470 --> 00:36:47.860

This technique keeps local state on the client

705

00:36:47.860 --> 00:36:52.270

and decides to retry or not according to the local state

706

00:36:52.270 --> 00:36:54.210

that it's keeping track of.

707

00:36:54.210 --> 00:36:56.260

The talking back targeting is being used

708

00:36:56.260 --> 00:36:58.220

to download this property.

709

00:36:58.220 --> 00:37:01.050

Let's say we have a client that's making a thousand

710

00:37:01.050 --> 00:37:04.630

requests per second, that's the purple line that I have

711

00:37:04.630 --> 00:37:06.217

on the graph.

712

00:37:06.217 --> 00:37:07.860

And it's static, it's not changing,

713

00:37:07.860 --> 00:37:10.850

always making a thousand requests per second.

714

00:37:10.850 --> 00:37:14.900

And the red line on the graph is the failure rate.

715

00:37:14.900 --> 00:37:18.080

Let's imagine that the service starting to have failure,

716

00:37:18.080 --> 00:37:21.357

it's initially it's healthy, it runs tests,

717

00:37:21.357 --> 00:37:24.083

having more and more failures, by the middle of the graph

718

00:37:24.083 --> 00:37:27.030

there's a hundred percent failures going on in the system.

719

00:37:27.030 --> 00:37:30.940

And then the failure start to get better over time.

720

00:37:30.940 --> 00:37:31.793

Back to zero.

721

00:37:32.800 --> 00:37:35.710

So let's look at the system when there is no

722

00:37:35.710 --> 00:37:37.990

client throttling is used.

723

00:37:37.990 --> 00:37:42.660

What happens is, the client starts retrying

724

00:37:42.660 --> 00:37:46.650

as it starts to observe the failures, and that excess

725

00:37:46.650 --> 00:37:49.330

traffic that's being created by the clients

726

00:37:49.330 --> 00:37:53.540

is following the same shape

727

00:37:53.540 --> 00:37:56.650

that the failure graph that we were just looking at, right?

728

00:37:56.650 --> 00:37:58.940

As there are more failures, there is more work

729

00:37:58.940 --> 00:38:00.440

being created.

730

00:38:00.440 --> 00:38:02.300

And towards the end off the middle of the graph,

731

00:38:02.300 --> 00:38:06.620

as the service is at hundred percent, let's say,

732

00:38:06.620 --> 00:38:11.560

the failure rate is the max, the retry count is the maximum,

733

00:38:11.560 --> 00:38:14.600

is the maximum amount of traffic being generated,

734

00:38:14.600 --> 00:38:16.483

and none of that being served.

735

00:38:17.410 --> 00:38:22.010

So, retrying in this case is not helping to resolve

736

00:38:22.010 --> 00:38:23.600

the issue, right?

737

00:38:23.600 --> 00:38:27.313

So let's see how client throttling helps this problem.

738

00:38:29.380 --> 00:38:32.223

So when the client throttling is enabled,

739

00:38:33.470 --> 00:38:36.840

as the service starts having more and more failures,

740

00:38:36.840 --> 00:38:39.797

the clients actually recognize it,

741

00:38:39.797 --> 00:38:43.930

and starts lowering the retry rate on the client side,

742

00:38:43.930 --> 00:38:46.310

and by the time service is having a hundred percent

743

00:38:46.310 --> 00:38:48.190

failing in the middle of the graph,

744

00:38:48.190 --> 00:38:49.840

there are no retries going on,

745

00:38:49.840 --> 00:38:51.980

it's just a flat number of requests

746

00:38:51.980 --> 00:38:54.890

that are still being sent, and when failure is seen,

747

00:38:54.890 --> 00:38:58.840

no retry, because it knows that retrying will not help.

748

00:38:58.840 --> 00:39:01.820

But as soon as service starts recovering in the second half

749

00:39:01.820 --> 00:39:06.070

of the graph, the retries start to pick up again,

750

00:39:06.070 --> 00:39:09.150

because now it knows that there's a chance a second retry

751

00:39:09.150 --> 00:39:11.693

might actually get a successful response.

752

00:39:13.080 --> 00:39:15.720

We found that this technique is much more effective

753

00:39:15.720 --> 00:39:19.650

to improve available depository of client applications

754

00:39:19.650 --> 00:39:23.460

while not generating any unnecessary work on the system.

755

00:39:23.460 --> 00:39:27.990

All right, that concludes our ten points today,

756

00:39:27.990 --> 00:39:29.340

and thank you for watching.

757

00:39:30.817 --> 00:39:34.067

(cheerful music plays)

758

00:39:39.664 --> 00:39:43.247

(cheerful music continues)